


```
In [1]: #import packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
import tensorflow as tf

#Load data
from sklearn.datasets import load_digits
import matplotlib.pyplot as plt
digits= load_digits()

#Digit Visualization
plt.matshow(digits.images[0])
plt.show()
```



```
In [2]: #making training and testing sets
from sklearn.model_selection import train_test_split
x_train , x_test, y_train , y_test = train_test_split(digits.data, digits.target,
```

Part 1 : Regularization


```
In [3]: #L2 penalty Logistic Regression
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import StratifiedKFold
skf = StratifiedKFold(n_splits=5)
skf.get_n_splits(x_train, y_train)
acc_list=[]
c_list=[]
alpha_lasso = [1e-10, 1e-9, 1e-8, 1e-7, 1e-6, 1e-5, 1e-4,1e-3,1e-2,0.1,0.2,0.3
,0.4,0.5,0.6,0.7,0.8,0.9, 1, 10, 100 ]
for i in alpha_lasso:
    acc_sum=0
    for train_index, test_index in skf.split(x_train, y_train):
        x_tra, x_tes = x_train[train_index], x_train[test_index]
        y_tra, y_tes = y_train[train_index], y_train[test_index]
        mod1 = LogisticRegression(penalty='l2',C=i)
        mod1.fit(x_tra, y_tra) # fit
        y_pred = mod1.predict(x_tes) # predict
        acc_log = round(mod1.score(x_tes, y_tes) * 100, 2)
        acc_sum += acc_log
    acc_list.append(acc_sum/5.0)
    c_list.append(i)
```

```
In [4]: #Dataframe for graph
df1 = pd.DataFrame({'c': np.round(np.log10(c_list),2), 'accuracy': acc_list})
df1
```

```
Out[4]:
```

	accuracy	c
0	72.306	-10.00
1	72.376	-9.00
2	72.376	-8.00
3	73.142	-7.00
4	78.232	-6.00
5	87.004	-5.00
6	91.874	-4.00
7	94.650	-3.00
8	96.388	-2.00
9	96.868	-1.00
10	96.452	-0.70
11	96.244	-0.52
12	96.038	-0.40
13	96.178	-0.30
14	96.174	-0.22
15	96.104	-0.15
16	96.104	-0.10
17	96.032	-0.05
18	95.962	0.00
19	95.200	1.00
20	94.156	2.00

```
In [5]: #Graph c values vs Accuracy
df1.plot('c','accuracy')
plt.title('Accuracy vs C values')
plt.xlabel('C')
plt.ylabel('Accuracy')
plt.show()
```




```
In [6]: #L1 penalty Logistic Regression
skf2 = StratifiedKFold(n_splits=5)
skf2.get_n_splits(x_train, y_train)
acc_list2=[]
c_list2=[]
alpha_lasso = [1e-10, 1e-9, 1e-8, 1e-7, 1e-6, 1e-5, 1e-4, 1e-3, 1e-2, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1, 10, 100 ]
for i in alpha_lasso:
    acc_sum2=0
    for train_index, test_index in skf.split(x_train, y_train):
        x_tra, x_tes = x_train[train_index], x_train[test_index]
        y_tra, y_tes = y_train[train_index], y_train[test_index]
        mod2 = LogisticRegression(penalty='l1',C=i)
        mod2.fit(x_tra, y_tra) # fit
        y_pred2 = mod2.predict(x_tes) # predict
        acc_log2 = round(mod2.score(x_tes, y_tes) * 100, 2)
        acc_sum2 += acc_log2
    acc_list2.append(acc_sum2/5.0)
    c_list2.append(i)
```

```
In [7]: #Dataframe for graph
df2 = pd.DataFrame({'c': np.round(np.log10(c_list2),2), 'accuracy': (acc_list2)})
df2
```

```
Out[7]:
```

	accuracy	c
0	10.508	-10.00
1	10.508	-9.00
2	10.508	-8.00
3	10.508	-7.00
4	10.508	-6.00
5	10.508	-5.00
6	10.508	-4.00
7	74.602	-3.00
8	92.560	-2.00
9	96.524	-1.00
10	96.314	-0.70
11	96.312	-0.52
12	95.966	-0.40
13	95.688	-0.30
14	95.548	-0.22
15	95.616	-0.15
16	95.550	-0.10
17	95.548	-0.05
18	95.618	0.00
19	95.270	1.00
20	94.922	2.00

```
In [8]: #Graph c values vs Accuracy
df2.plot('c','accuracy')
plt.title('Accuracy vs C values')
plt.xlabel('C')
plt.ylabel('Accuracy')
plt.show()
```



Part 2 : Neural Networks

#Q1

```

In [9]: #setting up tensorboard
tf.reset_default_graph()

from datetime import datetime
import os
import pathlib

t = datetime.utcnow().strftime("%Y%m%d%H%M%S")
log_dir = "tf_logs"
logd = "/tmp/{}/r{}/".format(log_dir, t)

from pathlib import Path
home = str(Path.home())

logdir = os.path.join(os.sep, home, logd)

if not os.path.exists(logdir):
    os.makedirs(logdir)

# TensorBoard Graph visualizer
import numpy as np
from IPython.display import clear_output, Image, display, HTML

def strip_consts(graph_def, max_const_size=32):
    """Strip large constant values from graph_def."""
    strip_def = tf.GraphDef()
    for n0 in graph_def.node:
        n = strip_def.node.add()
        n.MergeFrom(n0)
        if n.op == 'Const':
            tensor = n.attr['value'].tensor
            size = len(tensor.tensor_content)
            if size > max_const_size:
                tensor.tensor_content = "<stripped %d bytes>" % size
    return strip_def

def show_graph(graph_def, max_const_size=32):
    """Visualize TensorFlow graph."""
    if hasattr(graph_def, 'as_graph_def'):
        graph_def = graph_def.as_graph_def()
    strip_def = strip_consts(graph_def, max_const_size=max_const_size)
    code = """
    <script src="//cdnjs.cloudflare.com/ajax/libs/polymer/0.3.3/platform.js">
    <script>
        function load() {{
            document.getElementById("{id}").pbtxt = {data};
        }}
    </script>
    <link rel="import" href="https://tensorboard.appspot.com/tf-graph-basic.b
    <div style="height:600px">
        <tf-graph-basic id="{id}"></tf-graph-basic>
    </div>
    """.format(data=repr(str(strip_def)), id='graph'+str(np.random.rand()))

    iframe = """
    <iframe seamless style="width:1200px;height:620px;border:0" srcdoc="{>

```

```
"".format(code.replace("'", '&quot;'))  
display(HTML(iframe))
```



```

In [10]: # Define input placeholder
x = tf.placeholder(tf.float32, shape = [None, 64])

W = tf.Variable(tf.zeros([64, 10])) # number of weights
b = tf.Variable(tf.zeros([10])) # number of bias terms

#softmax regression
y_hat = tf.nn.softmax(tf.matmul(x, W) + b)

#output placeholder
y = tf.placeholder(tf.float32, [None, 10])

# Cross entropy
ce = tf.reduce_mean(-tf.reduce_sum( y* tf.log(y_hat), axis=1))

#optimizing function to minimize loss with learning rate = 0.001
train_step = tf.train.GradientDescentOptimizer(0.001).minimize(ce)

#One Hot Encoder for training set
from sklearn.preprocessing import OneHotEncoder
data = y_train
onehot_encoder = OneHotEncoder(sparse=False)
data2 = data.reshape(len(data), 1)
y_train_enc = onehot_encoder.fit_transform(data2)

#Repeating the same for test set
data = y_test
onehot_encoder = OneHotEncoder(sparse=False)
data2 = data.reshape(len(data), 1)
y_test_enc = onehot_encoder.fit_transform(data2)

#defining function to make batches of training data
def next_batch_fun(num, x, y):
    idx = np.arange(0, len(data))
    np.random.shuffle(idx)
    idx = idx[:num]
    x_shuffle = [x[i] for i in idx]
    y_shuffle = [y[i] for i in idx]
    return np.asarray(x_shuffle), np.asarray(y_shuffle)

#list variables for graph later
ta = []
va = []
e = []

#training taking batch_size=100, epochs=200
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    for epoch in range(200):
        avg_cost = 0.
        total_batch = int(digits.target.size/100)
        # Loop over all batches
        for i in range(total_batch):

```

```

batch_xs, batch_ys = next_batch_fun(100,x_train,y_train_enc)
# Fit training using batch data
_, c = sess.run([train_step, ce], feed_dict={x: batch_xs, y: batch_ys})

# Compute average Loss
avg_cost += c / total_batch
# To compute logs per 10 epoch step for graph later
if (epoch+1) % 10 == 0:
    #accuracy
    correct_prediction = tf.equal(tf.argmax(y,1), tf.argmax(y_hat,1))
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
    acc_train = accuracy.eval(feed_dict={x:x_train , y: y_train_enc})
    acc_val = accuracy.eval(feed_dict={x: x_test,y: y_test_enc})
    ta.append(acc_train)
    va.append(acc_val)
    e.append(epoch+1)
#To display logs per 50 epochs
if (epoch+1) % 50 == 0:
    print ("Epoch:", '%04d' % (epoch+1), "cost=", "{:.9f}".format(avg_cost))
    print ("Train accuracy:", np.round((acc_train)*100,2),"% ", "Val accuracy:", np.round((acc_val)*100,2),"% ")

print ("Optimization Finished!")

#Dataframes for graphs
df_ta = pd.DataFrame({'epoch': e , 'training accuracy': ta})
df_va = pd.DataFrame({'epoch': e , 'validation accuracy': va})

#Plotting Epoch vs Training Accuracy
df_ta.plot('epoch','training accuracy')
plt.title('Epoch vs. Training Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Training Accuracy')
plt.show()



#Plotting Epoch vs Validation Accuracy
df_va.plot('epoch','validation accuracy')
plt.title('Epoch vs. Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Validation Accuracy')
plt.show()

```

```

Epoch: 0050 cost= 0.163905337
Train accuracy: 95.89 %   Val accuracy: 93.61 %
Epoch: 0100 cost= 0.111125199
Train accuracy: 96.31 %   Val accuracy: 93.06 %
Epoch: 0150 cost= 0.076180501
Train accuracy: 96.52 %   Val accuracy: 93.33 %
Epoch: 0200 cost= 0.065715626
Train accuracy: 96.59 %   Val accuracy: 93.06 %
Optimization Finished!

```



```
In [11]: #Tensorboard Graph  
show_graph(tf.get_default_graph())
```



#Q2

```

In [16]: #Hyperparameters
n_hidden1 = 300
n_hidden2 = 200
n_hidden3 = 100
n_outputs = 10

#resetting the tensorboard graph
tf.reset_default_graph()

#Input and output placeholders
X = tf.placeholder(tf.float32, shape=[None, 64], name="X")
Y = tf.placeholder(tf.int64, shape=[None], name="Y")

#Defining function to write weights, biases and loss for each layer
def neuron_layer(X, n_neurons, keep_prob, name, activation=None):
    # X input to neuron
    # number of neurons for the layer
    # name of layer
    # pass in eventual activation function

    with tf.name_scope(name):
        n_inputs = int(X.get_shape()[1])

        # initialize weights to prevent vanishing / exploding gradients
        stddev = 2 / np.sqrt(n_inputs)
        init = tf.truncated_normal((n_inputs, n_neurons), stddev=stddev)

        # Initialize weights for the layer
        W = tf.Variable(init, name="weights")
        # biases
        b = tf.Variable(tf.zeros([n_neurons]), name="bias")

        drop_out = tf.nn.dropout(X, keep_prob)

        # Output from every neuron
        Z = tf.matmul(drop_out, W) + b
        if activation is not None:
            return activation(Z)
        else:
            return Z

#DNN for each layer along with tanh activation
with tf.name_scope("dnn"):
    hidden1 = neuron_layer(X, n_hidden1, 0.9, name="hidden1", activation=tf.nn.tanh)
    hidden2 = neuron_layer(hidden1, n_hidden2, 0.9, name="hidden2", activation=tf.nn.tanh)
    hidden3 = neuron_layer(hidden2, n_hidden3, 0.9, name="hidden3", activation=tf.nn.tanh)
    logits = neuron_layer(hidden3, n_outputs, 0.9, name="outputs")

#Softmax regression
with tf.name_scope("loss"):
    # logits are from the last output of the dnn
    xentropy = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=Y, logits=logits)
    ce = tf.reduce_mean(xentropy, name="loss")

```

```

learning_rate = 0.01

#Optimizer
with tf.name_scope("train"):
    optimizer = tf.train.GradientDescentOptimizer(learning_rate)
    train_step = optimizer.minimize(ce)

#Accuracy
with tf.name_scope("eval"):
    correct = tf.nn.in_top_k(logits, Y, 1)
    accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))

#Making batches
def next_batch_fun(num, x, y):
    idx = np.arange(0 , len(y))
    np.random.shuffle(idx)
    idx = idx[:num]
    x_shuffle = [x[i] for i in idx]
    y_shuffle = [y[i] for i in idx]
    return np.asarray(x_shuffle), np.asarray(y_shuffle)

#List for graphs
ta2 = []
va2 = []
e2 = []

init = tf.global_variables_initializer()
saver = tf.train.Saver()

#train
with tf.Session() as sess:
    init.run()
    for epoch in range(1000):
        for iteration in range(int(digits.target.size/100)):
            batch_xs, batch_ys = next_batch_fun(100,x_train,y_train)
            sess.run(train_step, feed_dict={X: batch_xs, Y: batch_ys})
        if (epoch+1)%50==0:
            acc_train = accuracy.eval(feed_dict={X: batch_xs, Y: batch_ys})
            acc_val = accuracy.eval(feed_dict={X: x_test,Y: y_test})
            ta2.append(acc_train)
            va2.append(acc_val)
            e2.append(epoch+1)
            print(epoch+1,"Train accuracy:",np.round((acc_train)*100,2),"% ", "Va

    save_path = saver.save(sess, "./my_model_final.ckpt") # save model


#Dataframes for graphs
df_ta2 = pd.DataFrame({'epoch': e2 , 'training accuracy': ta2})
df_va2 = pd.DataFrame({'epoch': e2 , 'validation accuracy': va2})


#Graphs for Epoch vs Training Accuracy
df_ta2.plot('epoch','training accuracy')
plt.title('Epoch vs. Training Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Training Accuracy')
plt.show()

```

```
#Graphs for Epoch vs Validation Accuracy
df_va2.plot('epoch','validation accuracy')
plt.title('Epoch vs. Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Validation Accuracy')
plt.show()
```

50	Train accuracy: 100.0 %	Val accuracy: 95.83 %
100	Train accuracy: 100.0 %	Val accuracy: 96.11 %
150	Train accuracy: 100.0 %	Val accuracy: 96.39 %
200	Train accuracy: 100.0 %	Val accuracy: 96.39 %
250	Train accuracy: 100.0 %	Val accuracy: 96.39 %
300	Train accuracy: 100.0 %	Val accuracy: 96.39 %
350	Train accuracy: 100.0 %	Val accuracy: 96.39 %
400	Train accuracy: 100.0 %	Val accuracy: 96.67 %
450	Train accuracy: 100.0 %	Val accuracy: 96.67 %
500	Train accuracy: 100.0 %	Val accuracy: 96.67 %
550	Train accuracy: 100.0 %	Val accuracy: 96.67 %
600	Train accuracy: 100.0 %	Val accuracy: 96.67 %
650	Train accuracy: 100.0 %	Val accuracy: 96.67 %
700	Train accuracy: 100.0 %	Val accuracy: 96.67 %
750	Train accuracy: 100.0 %	Val accuracy: 96.67 %
800	Train accuracy: 100.0 %	Val accuracy: 96.67 %
850	Train accuracy: 100.0 %	Val accuracy: 96.67 %
900	Train accuracy: 100.0 %	Val accuracy: 96.67 %
950	Train accuracy: 100.0 %	Val accuracy: 96.67 %
1000	Train accuracy: 100.0 %	Val accuracy: 96.67 %






```
In [13]: #Tensorboard Graph  
show_graph(tf.get_default_graph())
```



```
In [ ]:
```